

Splunk ITSI

ITSI Architecture & KV Store

1. ITSI Deployment Model

What is ITSI in Splunk?

Splunk IT Service Intelligence (ITSI) is an application that runs on top of **Splunk Enterprise**. It does **not ingest data by itself** — it **consumes data already indexed** in Splunk (metrics, events, logs).

Core Components of ITSI Architecture

1. Search Head (SH)

- Hosts:
 - ITSI UI
 - Services
 - KPIs
 - Correlation searches
 - Glass Tables
- Executes:
 - KPI base searches
 - Correlation searches
- Heavy dependency on:
 - Scheduler
 - KV Store

- If Search Head is unhealthy → ITSI becomes unusable
-

2. Indexers

- Store:
 - Metrics (_metrics)
 - Events & logs
 - KPI results

- Execute:
 - Distributed searches initiated by Search Head
-

3. KV Store (MongoDB-based)

- Stores **ITSI configuration & runtime state**
- Acts as ITSI's **metadata database**

(Deep dive covered later)

4. Scheduler

- Executes:
 - KPI base searches
 - Correlation searches
- Controls:
 - Search frequency
 - Resource allocation

Poor scheduler performance = delayed KPIs

Common ITSI Deployment Models

1. Single Instance (POC / Demo)

- All components on one server
- Not recommended for production

Use case:

- Learning
 - Proof of concept
-

2. Distributed Deployment (Production)

- Dedicated:
 - Search Head
 - Indexers
- KV Store runs on Search Head

Most common **enterprise ITSI setup**

3. Search Head Cluster + Indexer Cluster

- Used in:
 - Large enterprises
 - High availability environments
- KV Store replicated across SHC members

 KV Store consistency is critical in SHC

2. Dependencies (ITSI Perspective)

What Are Dependencies in ITSI?

Dependencies represent **relationships between services or components** that affect **health propagation**.

Example:

Database → Application → Web

Types of Dependencies

1. Service-to-Service Dependencies

- Parent service health depends on child service
- Used for:
 - Impact analysis
 - Root cause identification

2. Entity-Level Dependencies

- Entity health impacts KPIs
- Example:
 - Host entity impacts CPU KPI

Dependency Direction

Direction Meaning

Upstream Root cause source

Direction	Meaning
-----------	---------

Downstream Impacted service

Why Dependencies Matter

- Enables:
 - Health roll-up
 - Failure propagation
 - Root cause analysis
 - Without dependencies:
 - ITSI becomes a KPI dashboard only
-

Dependency Best Practices

- Keep dependency depth minimal
 - Avoid circular dependencies
 - Model **technical → business flow**
-

3. KV Store Role (Very Critical Topic)

What Is KV Store?

- Embedded **NoSQL database** (MongoDB)
 - Runs inside Splunk (splunkd)
 - Used by apps like:
 - ITSI
 - ES
 - SOAR (some components)
-

What KV Store Stores for ITSI

Stored in KV Store Not Stored

Services	Raw metrics
KPIs	Logs
Thresholds	Events

Stored in KV Store Not Stored

Entities Indexed data

Episodes

👉 KV Store stores “what ITSI knows” — not “what Splunk collects”

Why KV Store Is Critical for ITSI

- All ITSI objects live here
- Real-time KPI state stored here
- Episode tracking depends on KV Store

🔴 If KV Store goes down:

- Services disappear
 - KPIs stop updating
 - Glass tables break
-

KV Store Architecture (Conceptual)

- Based on MongoDB
- Each ITSI object stored as a document
- Collections used for:
 - itsi_services
 - itsi_kpis
 - itsi_entities
 - itsi_episodes

KV Store in Search Head Clustering

- KV Store replicated across SHC members
- One node acts as:
 - Primary
- Others:
 - Secondary replicas

Failure scenarios:

- KV Store sync lag → ITSI inconsistency
 - Split brain → corrupted ITSI objects
-

Common KV Store Issues in ITSI

- Disk I/O latency
 - Corrupted collections
 - Memory pressure
 - Large episode history
-

KV Store Best Practices

- Use SSD storage
- Monitor:
 - KV Store status
 - MongoDB processes
- Take regular backups
- Limit episode retention

Entities & Entity Types (ITSI)

1. What Is an Entity in ITSI?

An **Entity** represents a **real-world, monitorable component** in your IT environment.

Examples:

- A **server** (Linux / Windows)
- A **database**
- An **application**
- A **network device**
- A **Kubernetes pod**
- A **cloud resource**

👉 **Key idea:**

Entities are the bridge between raw data and services.

Without entities:

- KPIs become static
 - Services cannot scale
 - Dependency modeling becomes difficult
-

2. Entity Types

What Is an Entity Type?

An **Entity Type** is a **classification** that defines:

- What kind of object the entity is
- What attributes it has
- What metrics apply to it

Think of it like a **template or schema** for entities.

Common Built-in Entity Types

Entity Type	Example
Host	Linux / Windows server
Application	Payment Service
Database	Oracle / MySQL
Network Device	Router / Switch
Service	API Gateway

Custom Entity Types

You can create custom entity types for:

- Microservices
- Kubernetes workloads
- Cloud-native components
- Business-level objects

Example:

- payment_microservice
- aws_alb

- k8s_namespace
-

Why Entity Types Matter

- Define **metric applicability**
- Enable **consistent KPI reuse**
- Simplify large-scale onboarding

Same entity type = same KPIs = faster service modeling

3. Entity Creation

Entities can be created **manually or automatically** based on incoming data.

Methods of Entity Creation

1. Automatic Entity Creation

- Most common method
- Based on:
 - Incoming metrics
 - Field extraction rules
- ITSI automatically detects entities when data arrives

Example:

host=web01

→ ITSI creates entity **web01** of type **Host**

2. Manual Entity Creation

Used when:

- No data yet
- Pre-defining services
- Business entities

Steps:

- ITSI UI → Configuration → Entities
- Define entity name
- Assign entity type

- Add attributes
-

3. Lookup-Based Entity Creation

- Used for bulk onboarding
- CSV lookup maps fields → entities

Example lookup:

```
entity_name,entity_type,environment,app  
web01,host,prod,payment
```

Entity Attributes

Attributes describe the entity and help in:

- Filtering
- Tagging
- Service assignment

Common attributes:

- Hostname
 - IP address
 - Environment (Prod / QA)
 - Application name
 - Region
-

4. Tagging

What Is Tagging?

Tagging allows you to **logically group entities** using labels.

Tags are:

- Key-value pairs
 - Dynamic
 - Reusable
-

Examples of Tags

env:prod

tier:frontend

app:payment

region:apac

Why Tagging Is Important

- Dynamic service composition
- Simplifies entity selection
- Enables scalable modeling

Example:

Instead of adding 50 hosts manually → add tag app:payment

Tagging Use Cases

- Environment separation
 - Application grouping
 - SLA-based services
 - Team ownership
-

Tagging Best Practices

- Use consistent naming
 - Avoid over-tagging
 - Prefer business-meaningful tags
-

5. Metric Mapping

What Is Metric Mapping?

Metric Mapping defines **how incoming metrics are associated with entities**.

Without mapping:

- KPIs won't calculate correctly
 - Health scores remain unknown
-

Metric Flow in ITSI

Metrics → Entity → KPI → Service Health

Metric Mapping Components

- Metric name
- Entity identifier (field)
- Entity type

Example:

cpu.utilization → host=web01 → Host entity

How Metric Mapping Works

- ITSI examines incoming metric fields
 - Matches them to entity attributes
 - Assigns metrics to the correct entity
-

Example: Host CPU Metric

Metric:

cpu.utilization

Fields:

host=web01

Mapping:

- Entity Type: Host
- Entity Identifier: host

Result:

- CPU KPI calculated per host entity
-

Metric Mapping Challenges

- Inconsistent field names
 - Missing entity identifiers
 - Duplicate entity names
-

Best Practices for Metric Mapping

- Standardize metric field names

- Use unique identifiers
 - Validate mappings early
 - Prefer metrics (mstats) over events
-

Key Relationship Summary

Component	Purpose
Entity Type	Classification
Entity	Real-world object
Tags	Logical grouping
Metric Mapping	Data association

Service & Template Modeling (ITSI)

1. What Is a Service in ITSI?

In ITSI, a **Service** represents **what the business actually cares about**, not just infrastructure.

Examples:

- Online Banking
- Payment Gateway
- E-Commerce Checkout
- Core Banking System

👉 **Key idea:**

Services translate technical metrics into business health.

2. Types of Services

1. Technical Services

- Infrastructure-level
- Example:

- Web Server Service
- Database Service

Used by:

- NOC
 - SRE teams
-

2. Business Services

- End-to-end workflows
- Built on top of technical services

Example:

E-Commerce Checkout

```
|--- Web Tier  
|--- App Tier  
└--- Database Tier
```

Used by:

- Management
 - Business stakeholders
-

3. Building Business Services

Step-by-Step Service Creation

Step 1: Define Service Scope

- What does the service deliver?
- What is the business outcome?

Example:

“Process online payments successfully”

Step 2: Add Entities

Entities represent:

- Hosts
- Applications
- Databases

- Containers

Selection methods:

- Manual selection
 - Tag-based selection (recommended)
-

Step 3: Add KPIs

KPIs measure:

- Performance
- Availability
- Errors
- Capacity

Examples:

- Response Time
 - Error Rate
 - CPU Utilization
 - Throughput
-

Step 4: KPI Aggregation

How multiple KPIs combine:

- Average
 - Weighted average
 - Worst case
-

Step 5: Assign Weights

- Assign importance to KPIs
- Business KPIs usually have higher weight

Example:

KPI	Weight
Error Rate	40%
Response Time	30%

KPI	Weight
-----	--------

Availability	30%
--------------	-----

Step 6: Define Dependencies

Dependencies define **health propagation** between services.

4. Templates

What Is a Service Template?

A **template** is a **reusable blueprint** for creating services.

It includes:

- KPIs
 - Thresholds
 - Dependencies
 - Aggregation logic
-

Why Templates Are Important

- Faster onboarding
- Standardization
- Reduced configuration errors
- Easy scaling

Example:

One “Web Tier Template” → 50 web services

Types of Templates

Template Type Usage

Technical	Web, DB, App tiers
-----------	--------------------

Business	End-to-end services
----------	---------------------

Entity-based	Host/Application templates
--------------	----------------------------

Template Inheritance

- Child services inherit:
 - KPIs
 - Thresholds
 - Can override specific values
-

5. Dependencies (Very Important Concept)

What Are Dependencies in ITSI?

Dependencies define **how service health impacts other services**.

Example:

Database → Application → Web → Business Service

Types of Dependencies

1. Parent–Child Dependencies

- Parent service health depends on child service
 - Used for root cause analysis
-

2. Peer Dependencies

- Services depend on each other equally
 - Used rarely
-

Dependency Direction

Direction Meaning

Upstream Root cause

Downstream Impacted services

Health Propagation

- If a child service turns red:
 - Parent service degrades automatically
 - Weight defines impact severity
-

6. Dependency Best Practices

- Model **real technical flow**
 - Avoid circular dependencies
 - Keep depth manageable
 - Test failure scenarios
-

7. Common Modeling Mistakes

- Treating infra as business services
 - Too many KPIs in one service
 - No dependency modeling
 - Equal weights everywhere
-

Key Relationship Summary

Entities → KPIs → Services → Dependencies → Business Health

KPI Base Searches & Thresholding (ITSI)

1. What Is a KPI in ITSI?

A **Key Performance Indicator (KPI)** is a **measurable signal** that represents the **health of an entity or service**.

Examples:

- Response Time
- Error Rate
- CPU Utilization
- Availability
- Throughput

👉 **Key concept:**

KPIs convert raw data into health scores.

2. KPI Architecture Flow

Raw Data → Base Search → KPI Value → Threshold → Health Score → Service Health

If the base search is wrong, **everything above it is wrong.**

3. Base Search Design

What Is a Base Search?

A **base search** is an SPL query that:

- Retrieves raw data (metrics or events)
 - Produces a **numeric time-series value**
 - Feeds one or more KPIs
-

Base Search Design Principles

1. Efficiency First

- Runs continuously
- Consumes scheduler resources
- Must be optimized

Best practice:

- Use mstats for metrics
 - Avoid heavy joins and subsearches
-

2. One Base Search → Multiple KPIs

- Base search fetches raw data once
- KPIs derive values from it

Example:

```
| mstats avg(cpu.utilization) as cpu_avg  
where host=* span=1m
```

KPIs:

- CPU Average
 - CPU Peak
 - CPU Variance
-

3. Time Span Selection

- Typical span: 1m or 5m
 - Smaller span = higher resolution + more load
 - Larger span = lower accuracy
-

4. Metrics vs Events

Type	Recommendation
Metrics	Preferred
Logs	Use only if metrics unavailable

Reason:

- Metrics are faster
 - Lower storage
 - Designed for time-series KPIs
-

Base Search Output Requirements

A valid base search must output:

- _time
- Entity identifier (host, service, app)
- Numeric value

If any are missing → KPI fails.

Common Base Search Mistakes

- Returning multiple values per entity
 - Missing _time
 - Using non-numeric fields
 - Over-filtering data
-

4. Thresholding

Thresholds convert **KPI values → health states**.

Health states:

- Green (Healthy)

- Yellow (Warning)
 - Red (Critical)
-

Static Thresholds

What Are Static Thresholds?

Fixed threshold values defined manually.

Example:

Value	Health
-------	--------

< 70%	Green
-------	-------

70–85%	Yellow
--------	--------

> 85%	Red
-------	-----

When to Use Static Thresholds

- Well-known limits
- Infrastructure metrics
- Compliance metrics

Examples:

- Disk usage
 - CPU utilization
 - Memory usage
-

Advantages

- Simple
 - Predictable
 - Easy to explain
-

Disadvantages

- Not adaptive
- Can generate false positives
- Doesn't handle seasonality

Adaptive Thresholds

What Are Adaptive Thresholds?

Dynamic thresholds calculated using **historical data** and **machine learning**.

ITSI uses:

- Past behavior
 - Time-of-day patterns
 - Day-of-week trends
-

When to Use Adaptive Thresholds

- Transaction volumes
 - Response times
 - Error rates
 - Business KPIs
-

Advantages

- Fewer false alerts
 - Handles spikes & seasonality
 - Self-adjusting
-

Disadvantages

- Requires sufficient historical data
 - Less predictable initially
 - Harder to explain to business
-

Static vs Adaptive Thresholds

Aspect	Static	Adaptive
Setup	Simple	Moderate
Flexibility	Low	High
False Positives	Higher	Lower

Aspect	Static	Adaptive
--------	--------	----------

Business KPIs	✗	✓
---------------	---	---

5. KPI Aggregation & Health Calculation

- KPI values evaluated per entity
- Health scores aggregated at service level
- Weighted average used by default

Example:

Response Time (Red) + Error Rate (Yellow)

→ Service Health = Red

6. Best Practices

- Start with static → move to adaptive
- Validate base search before KPI creation
- Use consistent time spans
- Monitor scheduler load

Correlation Searches & Episode Review (ITSI)

1. Why Correlation & Episodes Matter

In real environments:

- Hundreds of alerts fire for **one incident**
- Operators waste time chasing symptoms

👉 ITSI's job here:

Group related events, identify the real issue, and show impact.

This is done using:

- **Correlation Searches**
- **Event Aggregation**
- **Episodes**
- **Root Cause Analysis (RCA)**

2. Correlation Searches

What Is a Correlation Search?

A **correlation search** is an SPL-based search that:

- Looks at **events, KPI alerts, or notable events**
 - Detects patterns across time, entities, or services
 - Generates **notable events** or feeds **episodes**
-

What Correlation Searches Can Use

- KPI threshold breaches
 - Log events
 - Metric anomalies
 - External alerts (APM, cloud, security)
-

Correlation Search Use Cases

- High CPU + High error rate
 - Multiple hosts failing in same service
 - Repeated KPI flapping
 - SLA breach patterns
-

Correlation Search Output

- Notable event
 - Episode contribution
 - Impacted service/entity
 - Severity & time window
-

Key Characteristics

- Scheduled (runs periodically)
 - Rule-based or pattern-based
 - Designed to **reduce alert noise**
-

Good vs Bad Correlation

Bad	Good
One alert per KPI	One episode per incident
No context	Business impact visible
Symptom-based	Cause-based

3. Event Aggregation

What Is Event Aggregation?

Event aggregation groups **multiple related events** into **one logical incident**, called an **Episode**.

Why Aggregation Is Needed

Without aggregation:

- 50 alerts = 50 tickets
- With aggregation:

- 50 alerts = **1 episode**
-

Aggregation Criteria

Events are grouped based on:

- Time window
- Entity
- Service
- Severity
- KPI type

Example:

CPU spike (Host A)

Memory spike (Host A)

Disk latency (Host A)

→ 1 Episode

Aggregation Methods

- Time-based grouping
 - Entity-based grouping
 - Service-based grouping
 - Rule-based grouping
-

Aggregation Window

Defines how long events are grouped together.

Example:

- 5 minutes
- 15 minutes
- 30 minutes

👉 Too small → too many episodes

👉 Too large → unrelated events grouped

4. Episodes

What Is an Episode?

An **Episode** represents:

- A **single incident**
 - With **multiple contributing events or KPIs**
 - Across entities and services
-

What an Episode Contains

- Start time & duration
 - Severity
 - Contributing KPIs
 - Impacted services
 - Suspected root cause
-

Episode Lifecycle

1. Episode opens
2. Events added

3. Root cause suggested
 4. Episode closes (manual or auto)
-

Episode Benefits

- Single pane of glass
 - Faster triage
 - Reduced MTTR
 - Clear business impact
-

5. Root Cause Analysis (RCA)

What Is RCA in ITSI?

RCA is ITSI's ability to:

- Identify **where the problem started**
 - Separate **cause vs symptom**
 - Show **impact propagation**
-

How ITSI Performs RCA

ITSI uses:

- Service dependencies
 - KPI health timelines
 - Event order
 - Entity relationships
-

RCA Example

Database latency ↑

→ App response time ↑

→ Web errors ↑

→ Checkout service ↓

Root Cause:

→ **Database latency**

RCA Visualization

- Dependency graphs
 - Health degradation timeline
 - Episode contributor ranking
-

Probable Root Cause

ITSI assigns:

- Confidence score
 - Based on earliest degradation
 - Dependency position (upstream)
-

6. Relationship Between Concepts

KPI Alerts

↓

Correlation Searches

↓

Event Aggregation

↓

Episode

↓

Root Cause Analysis

7. Best Practices

Correlation Searches

- Keep rules simple
 - Correlate only meaningful KPIs
 - Avoid duplicate rules
-

Event Aggregation

- Tune time windows carefully
- Group by entity/service

- Avoid over-aggregation
-

RCA

- Model dependencies correctly
 - Ensure KPI timestamps are accurate
 - Avoid circular dependencies
-

8. Common Mistakes

- No dependency modeling → bad RCA
 - Correlating everything → noise
 - Too many correlation searches
 - Ignoring episode tuning
-

Service Analyzer & Glass Tables (ITSI)

1. Service Analyzer

What Is Service Analyzer?

The **Service Analyzer** is ITSI's **primary operational view** that shows:

- Overall service health
- KPI health breakdown
- Dependency-driven impact
- Active episodes

👉 Think of it as the **mission control** for IT operations.

Purpose of Service Analyzer

- Provide **single-pane visibility**
- Help NOC/SRE teams:
 - Detect issues fast
 - Prioritize incidents
 - Understand business impact

What Service Analyzer Displays

- Service name
 - Health score (0–100)
 - Health color:
 - Green (Healthy)
 - Yellow (Degraded)
 - Red (Critical)
 - Child services & dependencies
 - Linked episodes
-

Service Hierarchy View

Services are displayed in a **tree structure**:

Business Service

|— Application Service

| |— Web Tier

| |— App Tier

|— Database Service

👉 Health **rolls up** from bottom → top.

Drilldowns in Service Analyzer

From a single click, users can:

- Open KPI details
 - View entity-level health
 - See contributing episodes
 - Jump to Glass Tables
-

2. Health Scores

What Is a Health Score?

A **numeric value (0–100)** representing service or KPI health.

Score Range Meaning

90–100 Healthy

70–89 Warning

< 70 Critical

How Health Scores Are Calculated

Health score is based on:

- KPI values
- Threshold breaches
- KPI weights
- Dependency impact

Example:

Error Rate (Red, weight 40%)

Response Time (Yellow, weight 30%)

Availability (Green, weight 30%)

→ Overall Health = Red

Why Health Scores Matter

- Abstracts raw metrics
- Enables **quick decisions**
- Makes IT data understandable for business users

Business leaders don't want CPU charts — they want "**Is my service OK?**"

3. Business Visuals

What Are Business Visuals?

Visual representations that map **technical health → business process flow**.

Used in:

- Executive reviews
- War rooms

- Incident calls
 - NOC displays
-

Business Visual Examples

- Order flow
 - Payment lifecycle
 - Login → Checkout journey
 - Transaction pipeline
-

4. Glass Tables

What Is a Glass Table?

A **custom, interactive visual dashboard** that overlays:

- Services
- KPIs
- Health scores
- Dependencies

On top of:

- Images
 - Diagrams
 - Architecture layouts
-

Why Glass Tables Are Powerful

- Visual storytelling
- Real-time health
- Business + IT alignment
- Extremely intuitive

👉 This is the **most demo-worthy ITSI feature**.

Glass Table Components

1. Shapes

- Represent services or entities

- Change color based on health
-

2. KPIs

- Embedded directly on visuals
 - Show live values and trends
-

3. Icons & Images

- Application icons
 - Architecture diagrams
 - Data flow illustrations
-

4. Interactions

- Click to drill down
 - Hover for KPI details
 - Navigate to Service Analyzer
-

Glass Table Health Behavior

- Colors auto-update in real time
- Health propagation follows dependencies
- Episodes can be highlighted visually

Example:

Database icon turns Red

→ App turns Yellow

→ Business Service turns Red

5. Real-Time Dashboards

What Makes Them “Real-Time”?

- Auto-refresh
 - KPI base search updates
 - Live episode integration
-

Real-Time Dashboard Use Cases

- NOC wallboards
 - 24x7 monitoring
 - Incident bridges
 - Executive status screens
-

Dashboard Characteristics

- Minimal clutter
 - Focus on health, not raw logs
 - Visual-first, data-second
-

Best Practices for Service Analyzer

- Limit KPI count per service
 - Use meaningful weights
 - Keep hierarchy simple
 - Regularly review stale services
-

Best Practices for Glass Tables

- Design for non-technical users
 - Avoid too many KPIs on one screen
 - Use consistent color semantics
 - Align visuals with real architecture
-

Common Mistakes

- Treating Glass Tables like dashboards
 - Overloading visuals
 - No drilldowns
 - Poor dependency modeling
-

Concept Relationship Summary

KPIs → Health Scores

Health Scores → Services

Services → Service Analyzer

Services → Glass Tables

Episodes → Visual Impact